

Employee Turn-Over EDA and Predictive Models

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from warnings import filterwarnings
filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv('Employee-turnover.csv')
df.head()
```

Out[2]:

	Employee Count	Employee ID	Department	Job Role	Gender	Age	Marital Status	Education	Education Field	Business Travel	...
0	1	1	Sales	Sales Executive	Female	41	Single	College	Life Sciences	Travel_Rarely	...
1	1	2	Research & Development	Research Scientist	Male	49	Married	Below College	Life Sciences	Travel_Frequently	...
2	1	4	Research & Development	Laboratory Technician	Male	37	Single	College	Other	Travel_Rarely	...
3	1	5	Research & Development	Research Scientist	Female	33	Married	Master	Life Sciences	Travel_Frequently	...
4	1	7	Research & Development	Laboratory Technician	Male	27	Married	Below College	Medical	Travel_Rarely	...

5 rows × 29 columns

```
In [3]: df.shape
```

Out[3]: (1470, 29)

```
In [4]: df.isnull().sum()
```

Out[4]:

Employee Count	0
Employee ID	0
Department	0
Job Role	0
Gender	0
Age	0
Marital Status	0
Education	0
Education Field	0
Business Travel	0
Distance From Home (kms)	0
Job Involvement	0
Job Level	0
Job Satisfaction	0
Monthly Income (USD)	0
Salary Hike (%)	0
Stock Option Level	0

```
Over Time                                0
No. of Companies Worked                  0
Total Working Years                      0
Years At Company                        0
Years In Current Role                   0
Years Since Last Promotion               0
Years With Curr Manager                  0
Environment Satisfaction                 0
Training Times Last Year                 0
Work Life Balance                       0
Relationship Satisfaction                 0
Attrition (Yes/No)                      0
dtype: int64
```

```
In [5]: df['Employee Count'].describe()
```

```
Out[5]: count      1470.0
mean         1.0
std          0.0
min          1.0
25%          1.0
50%          1.0
75%          1.0
max          1.0
Name: Employee Count, dtype: float64
```

```
In [6]: # We can remove this column since it doesn't help us in any way, values are same for all
# no deviation and doesn't make sense
df.drop('Employee Count',axis=1,inplace=True)
df
```

Out[6]:

	Employee ID	Department	Job Role	Gender	Age	Marital Status	Education	Education Field	Business Travel	Distance From Home (km)
0	1	Sales	Sales Executive	Female	41	Single	College	Life Sciences	Travel_Rarely	
1	2	Research & Development	Research Scientist	Male	49	Married	Below College	Life Sciences	Travel_Frequently	
2	4	Research & Development	Laboratory Technician	Male	37	Single	College	Other	Travel_Rarely	
3	5	Research & Development	Research Scientist	Female	33	Married	Master	Life Sciences	Travel_Frequently	
4	7	Research & Development	Laboratory Technician	Male	27	Married	Below College	Medical	Travel_Rarely	
...	
1465	2061	Research & Development	Laboratory Technician	Male	36	Married	College	Medical	Travel_Frequently	2
1466	2062	Research & Development	Healthcare Representative	Male	39	Married	Below College	Medical	Travel_Rarely	
1467	2064	Research & Development	Manufacturing Director	Male	27	Married	Bachelor	Life Sciences	Travel_Rarely	
1468	2065	Sales	Sales Executive	Male	49	Married	Bachelor	Medical	Travel_Frequently	

	Employee ID	Department	Job Role	Gender	Age	Marital Status	Education	Education Field	Business Travel	Distance From Home (kms)
1469	2068	Research & Development	Laboratory Technician	Male	34	Married	Bachelor	Medical	Travel_Rarely	

1470 rows × 28 columns

In [7]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 28 columns):
 #   Column                                          Non-Null Count  Dtype
---  -
 0   Employee ID                                  1470 non-null   int64
 1   Department                                  1470 non-null   object
 2   Job Role                                     1470 non-null   object
 3   Gender                                       1470 non-null   object
 4   Age                                          1470 non-null   int64
 5   Marital Status                             1470 non-null   object
 6   Education                                    1470 non-null   object
 7   Education Field                             1470 non-null   object
 8   Business Travel                             1470 non-null   object
 9   Distance From Home (kms)                   1470 non-null   int64
10   Job Involvement                             1470 non-null   object
11   Job Level                                   1470 non-null   int64
12   Job Satisfaction                           1470 non-null   object
13   Monthly Income (USD)                       1470 non-null   int64
14   Salary Hike (%)                            1470 non-null   int64
15   Stock Option Level                         1470 non-null   int64
16   Over Time                                  1470 non-null   object
17   No. of Companies Worked                     1470 non-null   int64
18   Total Working Years                         1470 non-null   int64
19   Years At Company                           1470 non-null   int64
20   Years In Current Role                      1470 non-null   int64
21   Years Since Last Promotion                 1470 non-null   int64
22   Years With Curr Manager                    1470 non-null   int64
23   Environment Satisfaction                   1470 non-null   object
24   Training Times Last Year                   1470 non-null   int64
25   Work Life Balance                          1470 non-null   object
26   Relationship Satisfaction                   1470 non-null   object
27   Attrition (Yes/No)                         1470 non-null   object
dtypes: int64(14), object(14)
memory usage: 321.7+ KB
```

In [8]:


```
df['Attrition'] = df['Attrition (Yes/No)'].apply(lambda x: 1 if x=='Yes' else 0)
df.head()
print(df['Attrition'].value_counts())
print(df['Attrition (Yes/No)'].value_counts())
```

```
0    1233
1     237
Name: Attrition, dtype: int64
No    1233
Yes    237
Name: Attrition (Yes/No), dtype: int64
```

In [9]:

```
#Department Column
plt.figure(figsize=(10,5))
```

```
df['Department'].value_counts(normalize=True).plot.pie(autopct='%1.f%%', shadow=True, explode=0.1)
plt.show()
```



1. 65% of employees are from R&D Department
2. 30% of employees are from Sales Dept.
3. ~5% of employees are from HR Dept.

```
In [10]: #Checking Attrition Rate for each category, we can create a function for that.
def calculate_attrition_rate(df,col):
    temp_df = pd.DataFrame(columns=['Category','Attrition Rate'])
    for cat in df[col].unique():
        attrition_count = df[(df[col]==cat) & (df['Attrition']==1)].shape[0]
        total_count = df[df[col]==cat].shape[0]
        temp_df = temp_df.append({'Category':cat, 'Attrition Rate':(attrition_count/total_count)})
    return temp_df.sort_values(by='Attrition Rate',ascending=False)
calculate_attrition_rate(df, 'Department')
```

```
Out[10]:
```

	Category	Attrition Rate
0	Sales	20.627803
2	Human Resources	19.047619
1	Research & Development	13.839750

```
In [11]: #Job Role
df['Job Role'].value_counts()
```


```
Out[11]:
```

Sales Executive	326
Research Scientist	292
Laboratory Technician	259
Manufacturing Director	145
Healthcare Representative	131
Manager	102
Sales Representative	83
Research Director	80
Human Resources	52

Name: Job Role, dtype: int64

```
In [12]: plt.figure(figsize=(10,8))
sns.countplot(x='Job Role',data=df,hue='Attrition')
```

```
plt.xticks(rotation=45)
plt.show()
```



```
In [13]: calculate_attrition_rate(df, 'Job Role')
```

```
Out[13]:
```

	Category	Attrition Rate
6	Sales Representative	39.759036
2	Laboratory Technician	23.938224
8	Human Resources	23.076923
0	Sales Executive	17.484663
1	Research Scientist	16.095890
3	Manufacturing Director	6.896552
4	Healthcare Representative	6.870229
5	Manager	4.901961
7	Research Director	2.500000

```
In [14]: #Check Gender only Male and Female
df['Gender'].value_counts()
```

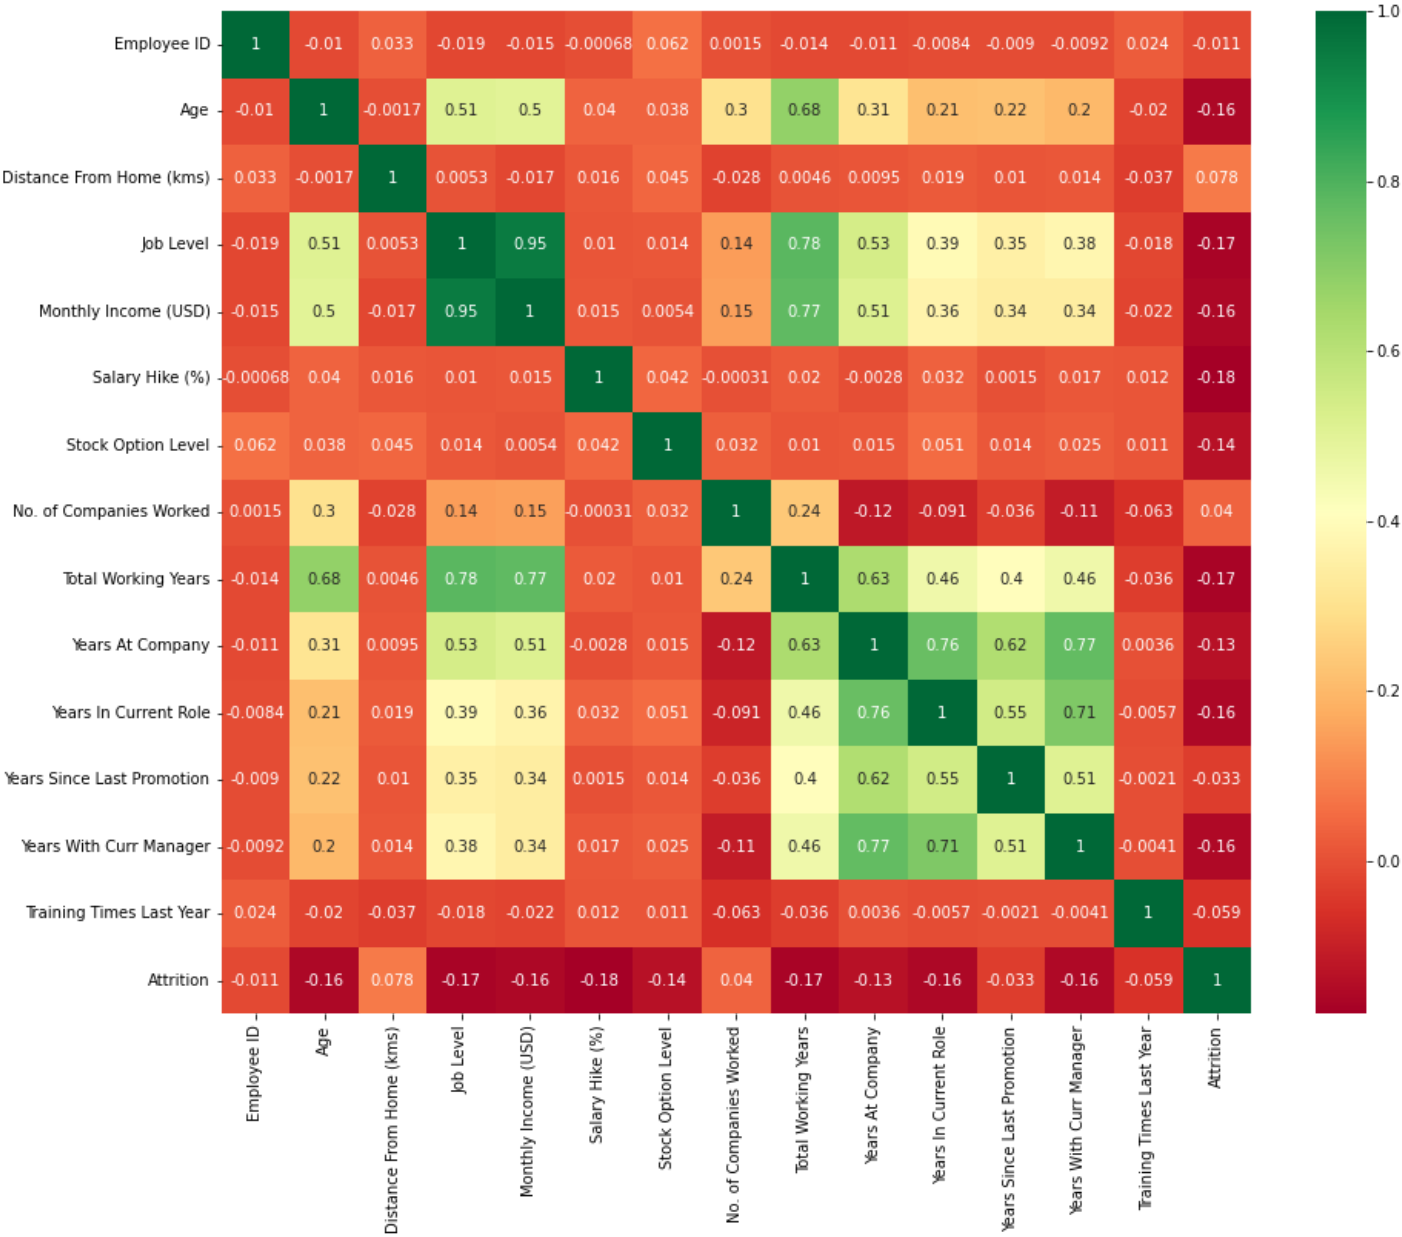
Out[14]: Male 882
Female 588
Name: Gender, dtype: int64

```
In [15]: calculate_attrition_rate(df, 'Gender')
```


Out[15]:

	Category	Attrition Rate
1	Male	17.006803
0	Female	14.795918

```
In [16]: corr = df.corr()  
plt.figure(figsize=(15,12))  
sns.heatmap(corr,annot=True,cmap='RdYlGn')  
plt.show()
```




```
In [17]: # We will study high correlations > 0.6 just to eliminate similar columns/ remove multi-c  
# For example, let's see the columns Monthly Income (USD) and Job Level with pearson coefi  
plt.figure(figsize=(10,5))  
sns.scatterplot(data=df,x='Job Level',y='Monthly Income (USD)',hue='Attrition')  
plt.show()
```



We can clearly see as Job Level rises, monthly income will rise as well. That's why having both these 2 columns for our model would be harmful hence we should remove 1 of them.

```
In [18]: sns.countplot(data=df, x='Job Level', hue='Attrition (Yes/No)')
```

```
Out[18]: <AxesSubplot:xlabel='Job Level', ylabel='count'>
```



As job level increases, number of attritions decreases.

We can also see this from Heatmap that Job level and Attrition have negative correlation.

```
In [19]: df.head()
```

```
Out[19]:
```

	Employee ID	Department	Job Role	Gender	Age	Marital Status	Education	Education Field	Business Travel	Distance From Home (kms)	...
0	1	Sales	Sales Executive	Female	41	Single	College	Life Sciences	Travel_Rarely	1	...
1	2	Research & Development	Research Scientist	Male	49	Married	Below College	Life Sciences	Travel_Frequently	8	...

	Employee ID	Department	Job Role	Gender	Age	Marital Status	Education	Education Field	Business Travel	Distance From Home (kms)	...
2	4	Research & Development	Laboratory Technician	Male	37	Single	College	Other	Travel_Rarely	2	...
3	5	Research & Development	Research Scientist	Female	33	Married	Master	Life Sciences	Travel_Frequently	3	...
4	7	Research & Development	Laboratory Technician	Male	27	Married	Below College	Medical	Travel_Rarely	2	...

5 rows × 29 columns

In [20]:

```
#We can either remove Job Level or Monthly Income (USD) column. I will remove income because it has a better correlation with other independent variables which is required.
#Even removing the total working years as it is highly correlated with job level.
df.drop(columns=['Monthly Income (USD)', 'Total Working Years'],axis=1,inplace=True)
df.head()
```

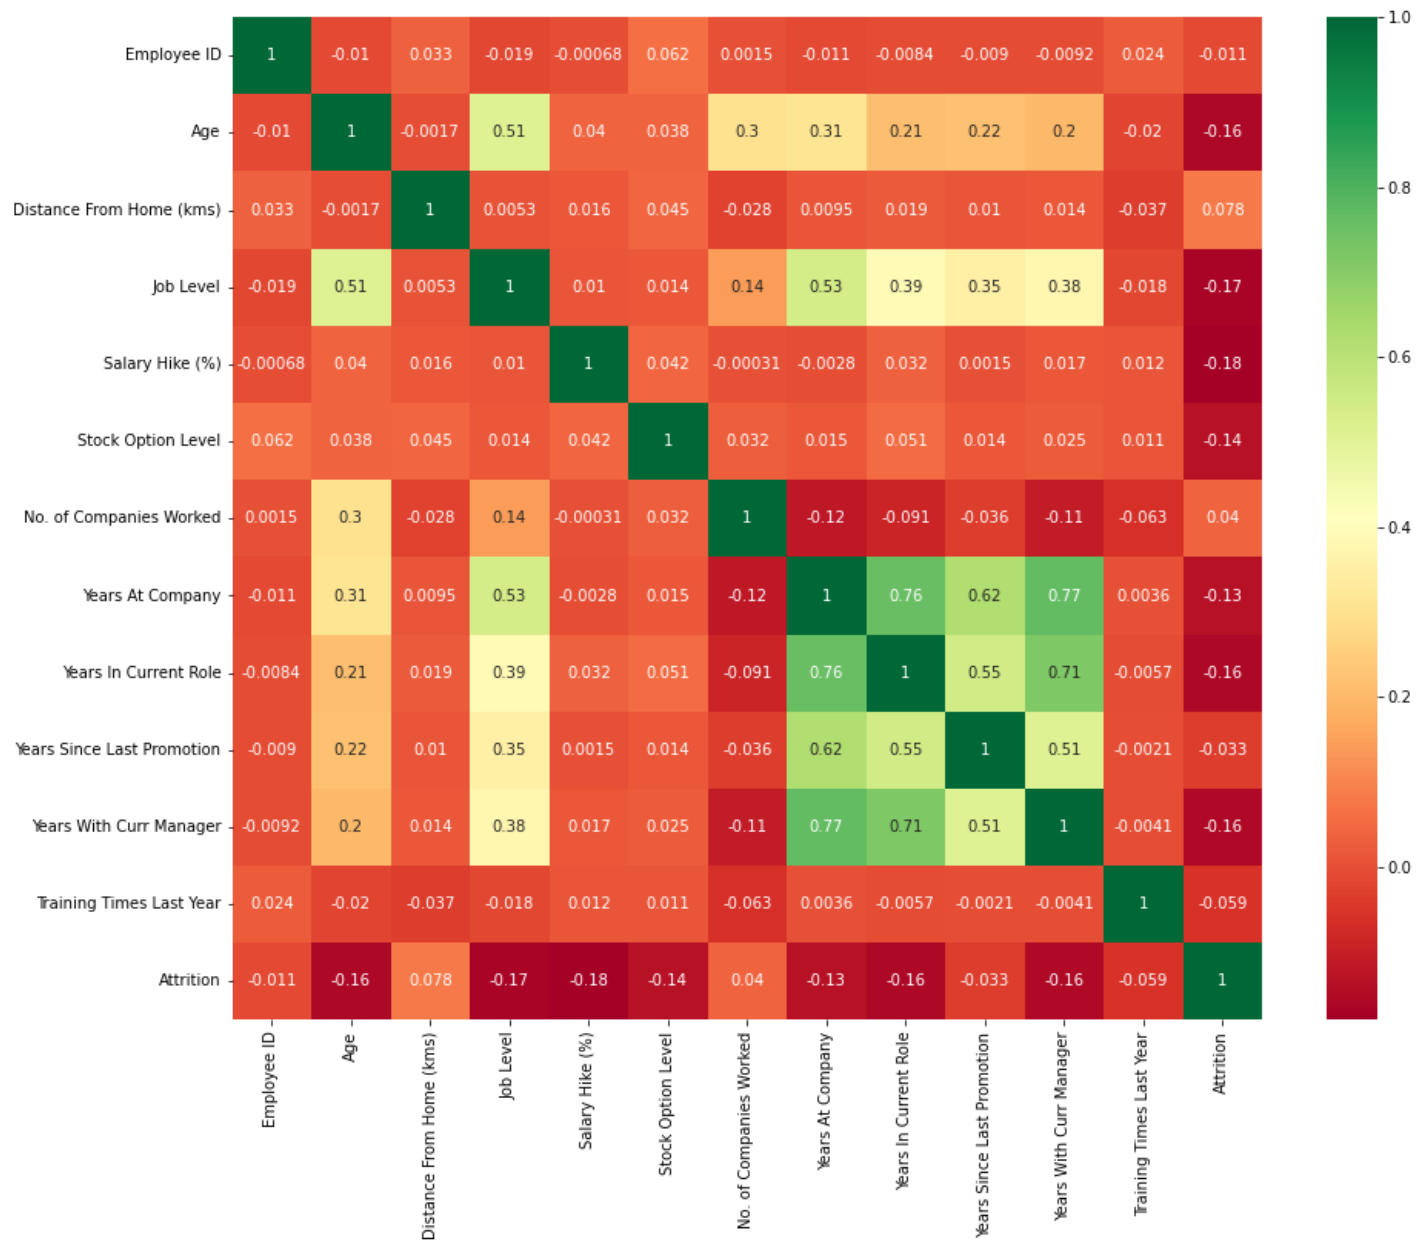
Out[20]:

	Employee ID	Department	Job Role	Gender	Age	Marital Status	Education	Education Field	Business Travel	Distance From Home (kms)	...
0	1	Sales	Sales Executive	Female	41	Single	College	Life Sciences	Travel_Rarely	1	...
1	2	Research & Development	Research Scientist	Male	49	Married	Below College	Life Sciences	Travel_Frequently	8	...
2	4	Research & Development	Laboratory Technician	Male	37	Single	College	Other	Travel_Rarely	2	...
3	5	Research & Development	Research Scientist	Female	33	Married	Master	Life Sciences	Travel_Frequently	3	...
4	7	Research & Development	Laboratory Technician	Male	27	Married	Below College	Medical	Travel_Rarely	2	...

5 rows × 27 columns

In [21]:

```
#Let's see the heatmap once again
corr = df.corr()
plt.figure(figsize=(15,12))
sns.heatmap(corr,annot=True,cmap='RdYlGn')
plt.show()
```

We see 'Years At Company' also has strong correlations with 'Years with Curr Manager' and 'Years in Current Role'. But I think these columns might be important so we won't remove them.

Also, we see that Attrition is most negatively correlated with Salary Hike %. It means, more the salary hike, less attrition.

Columns/Features which might be responsible for attrition:

1. Salary Hike %
2. Years with current manager
3. Years in Current Role
4. Years at a Company
5. Job Level
6. Stock Option Level
7. Age
8. Years Since Last Promotion


In [22]:

```
#let's try to understand Stock Option Level
df['Stock Option Level'].value_counts()
```

```
Out[22]: 0    631
          1    596
          2    158
          3     85
          Name: Stock Option Level, dtype: int64
```

```
In [23]: sns.countplot(data=df, x='Stock Option Level', hue='Attrition')
```

```
Out[23]: <AxesSubplot:xlabel='Stock Option Level', ylabel='count'>
```



```
In [24]: #Delete Attrition (Yes/No) column since we have our new Attrition column now.
df.drop(columns=['Attrition (Yes/No)'], inplace=True, axis=1)
```

```
In [25]: df.head()
```

```
Out[25]:
```

	Employee ID	Department	Job Role	Gender	Age	Marital Status	Education	Education Field	Business Travel	Distance From Home (kms)	...
0	1	Sales	Sales Executive	Female	41	Single	College	Life Sciences	Travel_Rarely	1	...
1	2	Research & Development	Research Scientist	Male	49	Married	Below College	Life Sciences	Travel_Frequently	8	...
2	4	Research & Development	Laboratory Technician	Male	37	Single	College	Other	Travel_Rarely	2	...
3	5	Research & Development	Research Scientist	Female	33	Married	Master	Life Sciences	Travel_Frequently	3	...
4	7	Research & Development	Laboratory Technician	Male	27	Married	Below College	Medical	Travel_Rarely	2	...

5 rows × 26 columns

Now, let's come to our categorical columns

```
In [26]: #Create a separate df for categorical columns
cat_df = df.select_dtypes(include='object').copy()
cat_df.head()
```

Out[26]:

	Department	Job Role	Gender	Marital Status	Education	Education Field	Business Travel	Job Involvement	Job Satisfaction	Ove Tim
0	Sales	Sales Executive	Female	Single	College	Life Sciences	Travel_Rarely	High	Very High	Ye
1	Research & Development	Research Scientist	Male	Married	Below College	Life Sciences	Travel_Frequently	Medium	Medium	Ni
2	Research & Development	Laboratory Technician	Male	Single	College	Other	Travel_Rarely	Medium	High	Ye
3	Research & Development	Research Scientist	Female	Married	Master	Life Sciences	Travel_Frequently	High	High	Ye
4	Research & Development	Laboratory Technician	Male	Married	Below College	Medical	Travel_Rarely	High	Medium	Ni

In [27]:

```
cat_df['Marital Status'].unique()
```

Out[27]:


array(['Single', 'Married', 'Divorced'], dtype=object)

In [28]:

```
sns.countplot(data=cat_df,x='Marital Status')
```

Out[28]:

<AxesSubplot:xlabel='Marital Status', ylabel='count'>



In [29]:


```
#Let's look at the relationship between Marital Status and Attrition
cat_df['Attrition'] = df.loc[:, 'Attrition'].copy()
cat_df.head()
```

Out[29]:

	Department	Job Role	Gender	Marital Status	Education	Education Field	Business Travel	Job Involvement	Job Satisfaction	Ove Tim
0	Sales	Sales Executive	Female	Single	College	Life Sciences	Travel_Rarely	High	Very High	Ye
1	Research & Development	Research Scientist	Male	Married	Below College	Life Sciences	Travel_Frequently	Medium	Medium	Ni
2	Research & Development	Laboratory Technician	Male	Single	College	Other	Travel_Rarely	Medium	High	Ye

	Department	Job Role	Gender	Marital Status	Education	Education Field	Business Travel	Job Involvement	Job Satisfaction	Over Time
3	Research & Development	Research Scientist	Female	Married	Master	Life Sciences	Travel_Frequently	High	High	Yes
4	Research & Development	Laboratory Technician	Male	Married	Below College	Medical	Travel_Rarely	High	Medium	No

```
In [30]: sns.countplot(data=cat_df, x='Marital Status', hue='Attrition')
plt.show()
```



```
In [31]: calculate_attrition_rate(cat_df, 'Marital Status')
```

```
Out[31]:
```

	Category	Attrition Rate
0	Single	25.531915
1	Married	12.481426
2	Divorced	10.091743

Single Employees have the most attrition rate.

This could be one of the driving factors for attrition rate.

```
In [32]: cat_df.columns
```

```
Out[32]: Index(['Department', 'Job Role', 'Gender', 'Marital Status', 'Education',
        'Education Field', 'Business Travel', 'Job Involvement',
        'Job Satisfaction', 'Over Time', 'Environment Satisfaction',
        'Work Life Balance', 'Relationship Satisfaction', 'Attrition'],
        dtype='object')
```


```
In [33]: #Education Field
cat_df['Education'].unique()
```

```
Out[33]: array(['College', 'Below College', 'Master', 'Bachelor', 'Doctor'],
        dtype=object)
```

```
In [34]: cat_df['Education'].value_counts()
```

```
Out[34]: Bachelor      572
Master      398
College     282
Below College 170
Doctor       48
Name: Education, dtype: int64
```

```
In [35]: sns.countplot(data=cat_df,x='Education',hue='Attrition')
plt.show()
```



```
In [36]: calculate_attrition_rate(cat_df,'Education')
```

```
Out[36]:
```

	Category	Attrition Rate
1	Below College	18.235294
3	Bachelor	17.307692
0	College	15.602837
2	Master	14.572864
4	Doctor	10.416667

Attrition Rate is almost similar among categories of different Education, hence we can drop this column.


```
In [37]: calculate_attrition_rate(cat_df,'Education Field')
```

```
Out[37]:
```

	Category	Attrition Rate
5	Human Resources	25.925926
4	Technical Degree	24.242424
3	Marketing	22.012579
0	Life Sciences	14.686469
2	Medical	13.577586
1	Other	13.414634


```
In [38]:
```

```
sns.countplot(data=cat_df,x='Education Field',hue='Attrition')
plt.show()
```



We can delete Education Field column as well

```
In [39]: #Checking Business Travel Column. This might be interesting.
sns.countplot(data=cat_df,x='Business Travel',hue='Attrition')
plt.show()
```



```
In [40]: calculate_attrition_rate(cat_df,'Business Travel')
```

```
Out[40]:
```

	Category	Attrition Rate
1	Travel_Frequently	24.909747
0	Travel_Rarely	14.956855
2	Non-Travel	8.000000

Huge Insight:


1. People who have to travel frequently have the highest attrition rate.
2. People who don't travel have the least attrition rate which is best for organization.

```
In [41]: cols_list = list(cat_df.columns)
```

cols_list

```
Out[41]: ['Department',
          'Job Role',
          'Gender',
          'Marital Status',
          'Education',
          'Education Field',
          'Business Travel',
          'Job Involvement',
          'Job Satisfaction',
          'Over Time',
          'Environment Satisfaction',
          'Work Life Balance',
          'Relationship Satisfaction',
          'Attrition']
```

```
In [42]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 6))
sns.countplot(data=cat_df, x='Job Involvement', hue='Attrition', ax=ax1)
sns.countplot(data=cat_df, x='Job Satisfaction', hue='Attrition', ax=ax2)
plt.show()
```



```
In [43]: calculate_attrition_rate(cat_df, 'Job Involvement')
```

```
Out[43]:
```

	Category	Attrition Rate
3	Low	33.734940
1	Medium	18.933333
0	High	14.400922
2	Very High	9.027778

```
In [44]: calculate_attrition_rate(cat_df, 'Job Satisfaction')
```

```
Out[44]:
```


	Category	Attrition Rate
--	----------	----------------

	Category	Attrition Rate
3	Low	22.837370
2	High	16.515837
1	Medium	16.428571
0	Very High	11.328976

Insights:

- 1. Employees like to get involved in their job. As the job involvement decreases, there is huge rise in attrition rate.
- 2. Employees with low job satisfaction have high attrition rate.

```
In [45]: plt.figure(figsize=(15,12))
sns.countplot(data=cat_df,x='Job Satisfaction',hue='Job Involvement')
plt.show()
```



```
In [46]: cat_df.head()
```

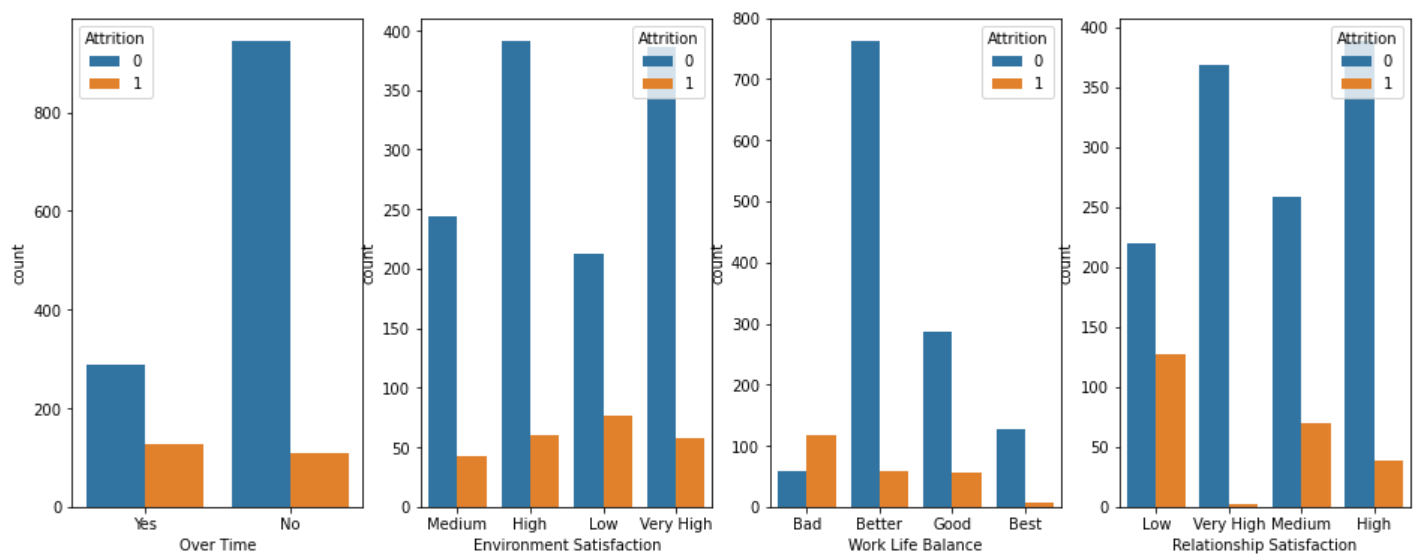
Out[46]:

Department	Job Role	Gender	Marital Status	Education	Education Field	Business Travel	Job Involvement	Job Satisfaction	Ove Tim
------------	----------	--------	----------------	-----------	-----------------	-----------------	-----------------	------------------	---------

	Department	Job Role	Gender	Marital Status	Education	Education Field	Business Travel	Job Involvement	Job Satisfaction	Over Time
0	Sales	Sales Executive	Female	Single	College	Life Sciences	Travel_Rarely	High	Very High	Yes
1	Research & Development	Research Scientist	Male	Married	Below College	Life Sciences	Travel_Frequently	Medium	Medium	No
2	Research & Development	Laboratory Technician	Male	Single	College	Other	Travel_Rarely	Medium	High	Yes
3	Research & Development	Research Scientist	Female	Married	Master	Life Sciences	Travel_Frequently	High	High	Yes
4	Research & Development	Laboratory Technician	Male	Married	Below College	Medical	Travel_Rarely	High	Medium	No

In [47]:

```
fig, (ax1,ax2,ax3,ax4) = plt.subplots(1,4,figsize=(16,6))
sns.countplot(data=cat_df,x='Over Time',hue='Attrition',ax=ax1)
sns.countplot(data=cat_df,x='Environment Satisfaction',hue='Attrition',ax=ax2)
sns.countplot(data=cat_df,x='Work Life Balance',hue='Attrition',ax=ax3)
sns.countplot(data=cat_df,x='Relationship Satisfaction',hue='Attrition',ax=ax4)
plt.show()
```



In [48]:

```
calculate_attrition_rate(cat_df,'Over Time')
```

Out[48]:

	Category	Attrition Rate
0	Yes	30.528846
1	No	10.436433

High Attrition Rate when employees have to do Over Time.

In [49]:

```
calculate_attrition_rate(cat_df,'Environment Satisfaction')
```

Out[49]:

	Category	Attrition Rate
2	Low	26.388889
0	Medium	14.982578

	Category	Attrition Rate
1	High	13.303769
3	Very High	13.063063

Employees don't want a very high satisfactory environment but it should be decent enough. A really bad environment could be a driver for attrition rates.

```
In [50]: calculate_attrition_rate(cat_df, 'Work Life Balance')
```

```
Out[50]:
```

	Category	Attrition Rate
0	Bad	66.666667
2	Good	16.129032
1	Better	7.073171
3	Best	4.545455

Huge Insight:

Bad Work Life Balance literally drives away the employee. Highest Attrition Rate in the data.

```
In [51]: calculate_attrition_rate(cat_df, 'Relationship Satisfaction')
```

```
Out[51]:
```

	Category	Attrition Rate
0	Low	36.705202
2	Medium	21.341463
3	High	8.920188
1	Very High	0.540541

Insights:

1. Low relationship satisfaction results in a high attrition rate.
2. Very High satisfaction has such low Attrition Rate.

Summary of Insights

1. Data is quite inconsistent because we have around 65% of employee from R&D Department, 30% from Sales and only 5% from HR Department.

2. We could clearly see as Job Level rises, monthly income will rise as well.

That's why having both these 2 columns for our model would be harmful hence we should remove 1 of them.

3. Single Employees have the most attrition rate.

This could be one of the driving factors for attrition rate.

4. Attrition Rate is almost similar among categories of different Education, hence we can drop this column. Similar for Education Field column.

5. People who have to travel frequently have the highest attrition rate

6. People who don't travel have the least attrition rate which is best for organization.

7. Employees like to get involved in their job. As the job involvement decreases, there is huge rise in

attrition rate.

8. Employees with low job satisfaction have high attrition rate.

9. High Attrition Rate when employees have to do Over Time.

10. Employees don't want a very high satisfactory environment but it should be decent enough. A really bad environment could be a driver for attrition rates.

11. Bad Work Life Balance literally drives away the employee. Highest Attrition Rate in the data.

12. Low relationship satisfaction results in a high attrition rate.

Very High satisfaction has such low Attrition Rate.

Key Driver Factors for Modeling

1. Salary Hike %
2. Years with current manager
3. Years in Current Role
4. Years at a Company
5. Job Level
6. Stock Option Level
7. Age
8. Marital Status
9. Business Travel
10. Job Satisfaction
11. Job Involvement
12. Over Time
13. Environment Satisfaction
14. Work Life Balance
15. Relationship Satisfaction
16. Department
17. Job Role
18. Gender
19. Years Since Last Promotion

Data Pre-Processing

```
In [52]: df.head()
```

Out[52]:

	Employee ID	Department	Job Role	Gender	Age	Marital Status	Education	Education Field	Business Travel	Distance From Home (kms)	...
0	1	Sales	Sales Executive	Female	41	Single	College	Life Sciences	Travel_Rarely	1	...
1	2	Research & Development	Research Scientist	Male	49	Married	Below College	Life Sciences	Travel_Frequently	8	...
2	4	Research & Development	Laboratory Technician	Male	37	Single	College	Other	Travel_Rarely	2	...
3	5	Research & Development	Research Scientist	Female	33	Married	Master	Life Sciences	Travel_Frequently	3	...

	Employee ID	Department	Job Role	Gender	Age	Marital Status	Education	Education Field	Business Travel	Distance From Home (kms)	...
4	7	Research & Development	Laboratory Technician	Male	27	Married	Below College	Medical	Travel_Rarely	2	...

5 rows × 26 columns

In [53]:

```
df.select_dtypes(include='object')
```

Out[53]:

	Department	Job Role	Gender	Marital Status	Education	Education Field	Business Travel	Job Involvement	Job Satisfaction
0	Sales	Sales Executive	Female	Single	College	Life Sciences	Travel_Rarely	High	Very High
1	Research & Development	Research Scientist	Male	Married	Below College	Life Sciences	Travel_Frequently	Medium	Medium
2	Research & Development	Laboratory Technician	Male	Single	College	Other	Travel_Rarely	Medium	High
3	Research & Development	Research Scientist	Female	Married	Master	Life Sciences	Travel_Frequently	High	High
4	Research & Development	Laboratory Technician	Male	Married	Below College	Medical	Travel_Rarely	High	Medium
...
1465	Research & Development	Laboratory Technician	Male	Married	College	Medical	Travel_Frequently	Very High	Very High
1466	Research & Development	Healthcare Representative	Male	Married	Below College	Medical	Travel_Rarely	Medium	Low
1467	Research & Development	Manufacturing Director	Male	Married	Bachelor	Life Sciences	Travel_Rarely	Very High	Medium
1468	Sales	Sales Executive	Male	Married	Bachelor	Medical	Travel_Frequently	Medium	Medium
1469	Research & Development	Laboratory Technician	Male	Married	Bachelor	Medical	Travel_Rarely	Very High	High

1470 rows × 13 columns

In [54]:

```
df.select_dtypes(include='int')
```

Out[54]:

	Employee ID	Age	Distance From Home (kms)	Job Level	Salary Hike (%)	Stock Option Level	No. of Companies Worked	Years At Company	Years In Current Role	Years Since Last Promotion	Years With Current Manager	Training Hours (Y)
0	1	41	1	2	11	0	2	6	4	0	5	
1	2	49	8	2	23	1	1	10	7	1	7	
2	4	37	2	1	15	0	6	0	0	0	0	
3	5	33	3	1	11	0	1	8	7	3	0	

	Employee ID	Age	Distance From Home (kms)	Job Level	Salary Hike (%)	Stock Option Level	No. of Companies Worked	Years At Company	Years In Current Role	Years Since Last Promotion	Years With Curr Manager	Train
	4	7	27	2	1	12	1	9	2	2	2	
	
	1465	2061	36	23	2	17	1	4	5	2	0	3
	1466	2062	39	6	3	15	1	4	7	7	1	7
	1467	2064	27	4	2	20	1	1	6	2	0	3
	1468	2065	49	2	2	14	0	2	9	6	0	8
	1469	2068	34	8	2	12	0	2	4	3	1	2

1470 rows × 13 columns

```
In [55]: ds = df.loc[:,['Department','Job Role','Gender','Marital Status','Business Travel','Job Involvement','Job Satisfaction','Over Time','Environment Satisfaction','Work Life Balance','Age','Job Level','Salary Hike (%)','Stock Option Level','Years At Company','Years With Curr Manager','Attrition']].copy()
ds.head()
```

```
Out[55]:
```

	Department	Job Role	Gender	Marital Status	Business Travel	Job Involvement	Job Satisfaction	Over Time	Environment Satisfaction	Work Life Balance
0	Sales	Sales Executive	Female	Single	Travel_Rarely	High	Very High	Yes	Medium	Best
1	Research & Development	Research Scientist	Male	Married	Travel_Frequently	Medium	Medium	No	High	Best
2	Research & Development	Laboratory Technician	Male	Single	Travel_Rarely	Medium	High	Yes	Low	Best
3	Research & Development	Research Scientist	Female	Married	Travel_Frequently	High	High	Yes	Very High	Best
4	Research & Development	Laboratory Technician	Male	Married	Travel_Rarely	High	Medium	No	Low	Best

```
In [56]: ds.shape
```

```
Out[56]: (1470, 20)
```

```
In [57]: #We have reduced the number of columns to 20. Great!
```

Setting Dependent and Independent variables

```
In [58]: X = ds.iloc[:, :-1].values
y=ds.iloc[:, -1].values
print(X)
print(y)
```

```
[[ 'Sales' 'Sales Executive' 'Female' ... 4 0 5]
```

```

['Research & Development' 'Research Scientist' 'Male' ... 7 1 7]
['Research & Development' 'Laboratory Technician' 'Male' ... 0 0 0]
...
['Research & Development' 'Manufacturing Director' 'Male' ... 2 0 3]
['Sales' 'Sales Executive' 'Male' ... 6 0 8]
['Research & Development' 'Laboratory Technician' 'Male' ... 3 1 2]]
[1 0 1 ... 0 0 0]

```

Encoding the Independent Variables

```

In [59]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder

```

```

In [60]: ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [0,1,2,3,4,5,6,7,8,9,10])]
X = np.array(ct.fit_transform(X))

```

```

In [61]: print(X)

[[0.0 0.0 1.0 ... 4 0 5]
 [0.0 1.0 0.0 ... 7 1 7]
 [0.0 1.0 0.0 ... 0 0 0]
 ...
 [0.0 1.0 0.0 ... 2 0 3]
 [0.0 0.0 1.0 ... 6 0 8]
 [0.0 1.0 0.0 ... 3 1 2]]

```

```

In [62]: X.shape

```

```

Out[62]: (1470, 50)

```

Splitting dataset into training and test set

```

In [63]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,random_state=0)

```

```

In [64]: print(X_train[0,42:])
#print(X_train)

[58 4 13 1 1 0 0 0]

```

```

In [65]: X_train.shape

```

```

Out[65]: (1176, 50)

```

```

In [66]: print(X_test)

[[0.0 0.0 1.0 ... 3 9 7]
 [0.0 1.0 0.0 ... 3 0 2]
 [0.0 0.0 1.0 ... 2 3 2]
 ...
 [0.0 0.0 1.0 ... 7 7 7]
 [0.0 1.0 0.0 ... 0 0 0]
 [0.0 1.0 0.0 ... 0 0 0]]

```

```

In [67]: print(y_train)

```

```
[0 0 0 ... 0 0 0]
```

```
In [68]: print(y_test)
```

```
[0 0 1 0 1 0 1 0 0 1 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0
 0 1 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 0 1 0 0 0 0 0 0 0 0 1 1 0 1 0 0 0 0
 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 1 0 0
 1 0 0 0 0 1 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 0 0 0 1 0
 0 0 0 1 0 0 1 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0
 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0
 1 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0]
```

Feature Scaling

```
In [69]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
#Gives value between -3 to +3
#We do not standardize our dummy variables.
X_train[:,42:] = sc.fit_transform(X_train[:,42:])
X_test[:,42:] = sc.transform(X_test[:,42:])
```

```
In [70]: print(X_train[0,42:])
```

```
[2.3389367036100883  1.759027295888783 -0.5394871849927129
 0.22505569171608666 -0.9820078512815037 -1.1568405801129382
-0.6738150570956313 -1.1504388612192216]
```

```
In [71]: print(X_test[0,42:])
```

```
[-0.08886551888941858  0.8545165394544985 -0.256887216372922
-0.9306888212014588  0.49540097572111685 -0.3303250745158451
 2.132819560889478  0.8143770527966774]
```

Machine Learning Classification Models

1. Logistic Regression
2. KNN
3. SVM
4. Kernel SVM
5. Naive Bayes
6. Decision Tree Classification
7. Random Forest Classification

Logistic Regression

```
In [72]: from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state=0)
classifier.fit(X_train,y_train)
```

```
Out[72]: LogisticRegression(random_state=0)
```

```
In [73]: ### Predicting test results
y_pred = classifier.predict(X_test)
```

```
print(np.concatenate((y_pred.reshape(len(y_pred),1),y_test.reshape(len(y_test),1)),1))
```

[illegible]

[0 0]
[0 0]
[1 1]
[1 1]
[0 0]
[0 0]
[1 0]
[1 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 1]
[1 0]
[0 0]
[0 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 1]
[0 0]
[0 0]
[0 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 1]
[0 0]
[1 0]
[0 1]
[0 0]
[0 0]
[0 0]


```

[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[1 1]
[0 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[1 0]
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 1]
[0 0]
[0 0]
[0 0]
[0 1]
[0 0]
[0 0]

```

In [74]:

```

#### Confusion Matrix
from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)

```

```

[[240   5]
 [ 23  26]]
0.9047619047619048

```

Out[74]:

Logistic Regression Model - 90.5% accuracy.

KNN Classifier

In [75]:

```

from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors=7, metric='minkowski', p=2)
classifier.fit(X_train, y_train)
#Predicting test result
y_pred = classifier.predict(X_test)
#print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.reshape(len(y_test),1)),1))
#### Confusion Matrix
from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)

```

```

[[244   1]
 [ 36  13]]

```

Out[75]: 0.8741496598639455

KNN Classifier Model - 87.4% accuracy.

SVM Classifier

In [76]:

```
from sklearn.svm import SVC
classifier = SVC(kernel = 'rbf', random_state=0)
classifier.fit(X_train,y_train)
#Predicting test result
y_pred = classifier.predict(X_test)
#print(np.concatenate((y_pred.reshape(len(y_pred),1),y_test.reshape(len(y_test),1)),1))
#### Confusion Matrix
from sklearn.metrics import confusion_matrix,accuracy_score
cm = confusion_matrix(y_test,y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[241   4]
 [ 30  19]]
0.8843537414965986
```

Out[76]:

SVM Classifier Model - 88.4% accuracy.

Decision Tree Classification

In [77]:

```
from sklearn.tree import DecisionTreeClassifier
classifier = DecisionTreeClassifier(criterion='entropy', random_state=0)
classifier.fit(X_train,y_train)
#Predicting test result
y_pred = classifier.predict(X_test)
#print(np.concatenate((y_pred.reshape(len(y_pred),1),y_test.reshape(len(y_test),1)),1))
#### Confusion Matrix
from sklearn.metrics import confusion_matrix,accuracy_score
cm = confusion_matrix(y_test,y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[224  21]
 [ 22  27]]
0.8537414965986394
```

Out[77]:

Decision Tree Classifier Model - 85.3% accuracy.

In [78]:

```
# Random Forest Classification
from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(n_estimators = 50,random_state=0)
classifier.fit(X_train,y_train)
#Predicting test result
y_pred = classifier.predict(X_test)
#print(np.concatenate((y_pred.reshape(len(y_pred),1),y_test.reshape(len(y_test),1)),1))
#### Confusion Matrix
from sklearn.metrics import confusion_matrix,accuracy_score
cm = confusion_matrix(y_test,y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[241   4]
 [ 30  19]]
```

Out[78]: 0.8843537414965986

Random Forest Classifier Model - 88.4% accuracy.

Result

We should use Logistic Regression Model (90.5% accuracy) since it gives us more than 90% accurate results.